


*Operating Systems:
Internals and Design Principles, 6/E*
William Stallings


Chapter 7


Memory Management



Roadmap


- Basic requirements of Memory Management
- Memory Partitioning
- Basic blocks of memory management
 - Paging
 - Segmentation






The need for memory management

- Memory is cheap today, and getting cheaper
 - But applications are demanding more and more memory, there is never enough!
- Memory Management, involves swapping blocks of data from secondary storage.
- Memory I/O is slow compared to a CPU
 - The OS must cleverly time the swapping to maximise the CPU's efficiency






Memory Management

Memory needs to be allocated to ensure a reasonable supply of ready processes to consume available processor time






Memory Management Requirements


- Relocation
- Protection
- Sharing
- Logical organisation
- Physical organisation






Requirements: Relocation

- The programmer does not know where the program will be placed in memory when it is executed,
 - it may be swapped to disk and return to main memory at a different location (relocated)
- Memory references must be translated to the actual physical memory address






Memory Management Terms

Table 7.1 Memory Management Terms

Term	Description
Frame	<i>Fixed</i> -length block of main memory.
Page	<i>Fixed</i> -length block of data in secondary memory (e.g. on disk).
Segment	<i>Variable-length</i> block of data that resides in secondary memory.



Addressing

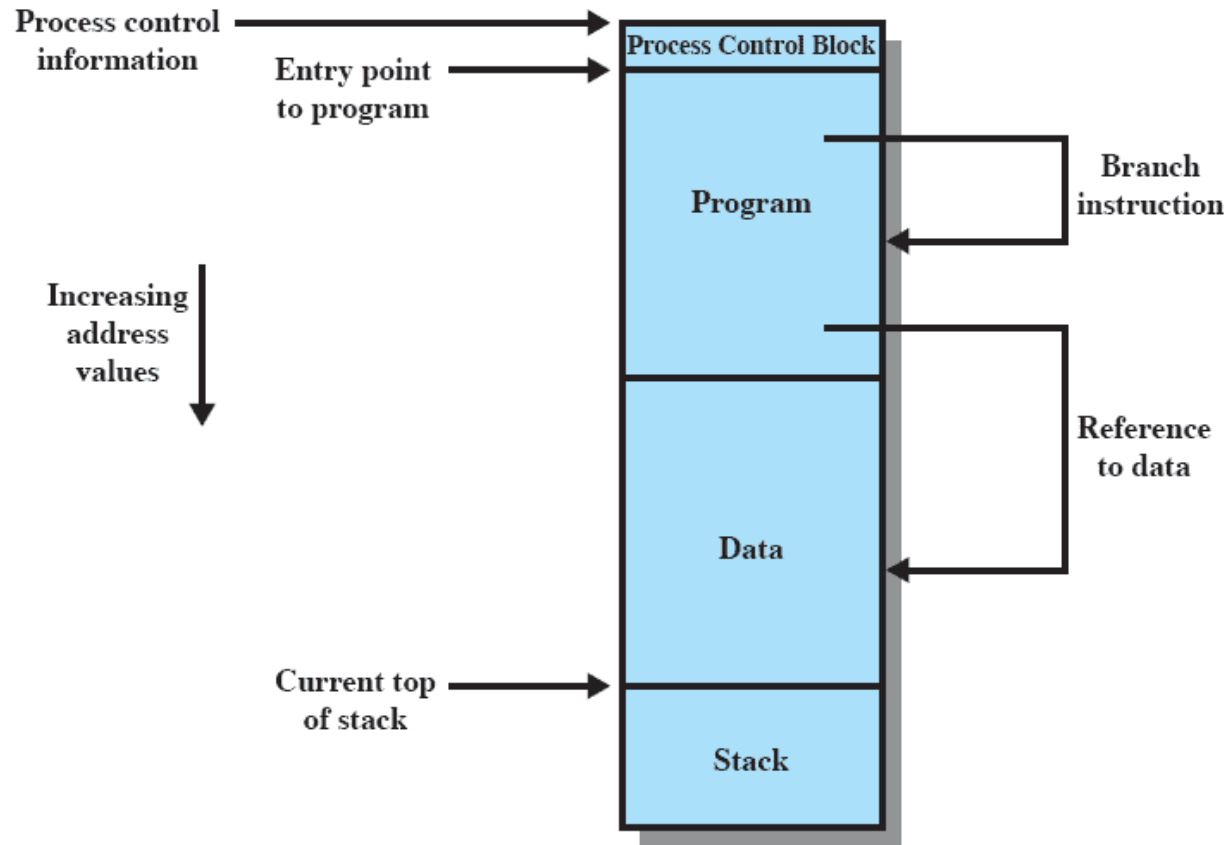



Figure 7.1 Addressing Requirements for a Process



Requirements: Protection

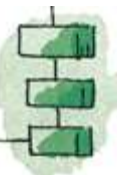
- Processes should not be able to reference memory locations in another process without permission
- Impossible to check absolute addresses at compile time
- Must be checked at run time






Requirements: Sharing

- Allow several processes to access the same portion of memory
- Better to allow each process access to the same copy of the program rather than have their own separate copy






Requirements: Logical Organization


- Memory is organized linearly (usually)
- Programs are written in modules
 - Modules can be written and compiled independently
- Different degrees of protection given to modules (read-only, execute-only)
- Share modules among processes
- Segmentation helps here





Requirements: Physical Organization

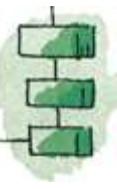
- Cannot leave the programmer with the responsibility to manage memory
- Memory available for a program plus its data may be insufficient
 - Overlaying allows various modules to be assigned the same region of memory but is time consuming to program
- Programmer does not know how much space will be available






Partitioning

- An early method of managing memory
 - Pre-virtual memory
 - Not used much now
- But, it will clarify the later discussion of virtual memory if we look first at partitioning
 - Virtual Memory has evolved from the partitioning methods






Types of Partitioning


- Fixed Partitioning
- Dynamic Partitioning
- Simple Paging
- Simple Segmentation
- Virtual Memory Paging
- Virtual Memory Segmentation






Fixed Partitioning

- Equal-size partitions (see fig 7.3a)
 - Any process whose size is less than or equal to the partition size can be loaded into an available partition
- The operating system can swap a process out of a partition
 - If none are in a ready or running state



(a) Equal-size partitions






Fixed Partitioning Problems

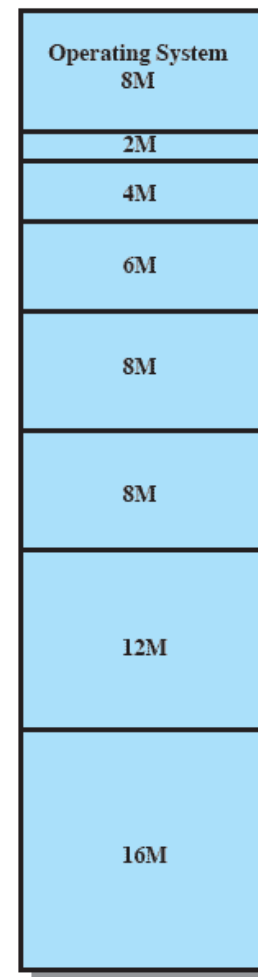
- A program may not fit in a partition.
 - The programmer must design the program with overlays
- Main memory use is inefficient.
 - Any program, no matter how small, occupies an entire partition.
 - This results in ***internal fragmentation***.






Solution – Unequal Size Partitions

- Lessens both problems
 - but doesn't solve completely
- In Fig 7.3b,
 - Programs up to 16M can be accommodated without overlay
 - Smaller programs can be placed in smaller partitions, reducing internal fragmentation




(b) Unequal-size partitions





Placement Algorithm

- Equal-size
 - Placement is trivial (no options)
- Unequal-size
 - Can assign each process to the smallest partition within which it will fit
 - Queue for each partition
 - Processes are assigned in such a way as to minimize wasted memory within a partition



Fixed Partitioning

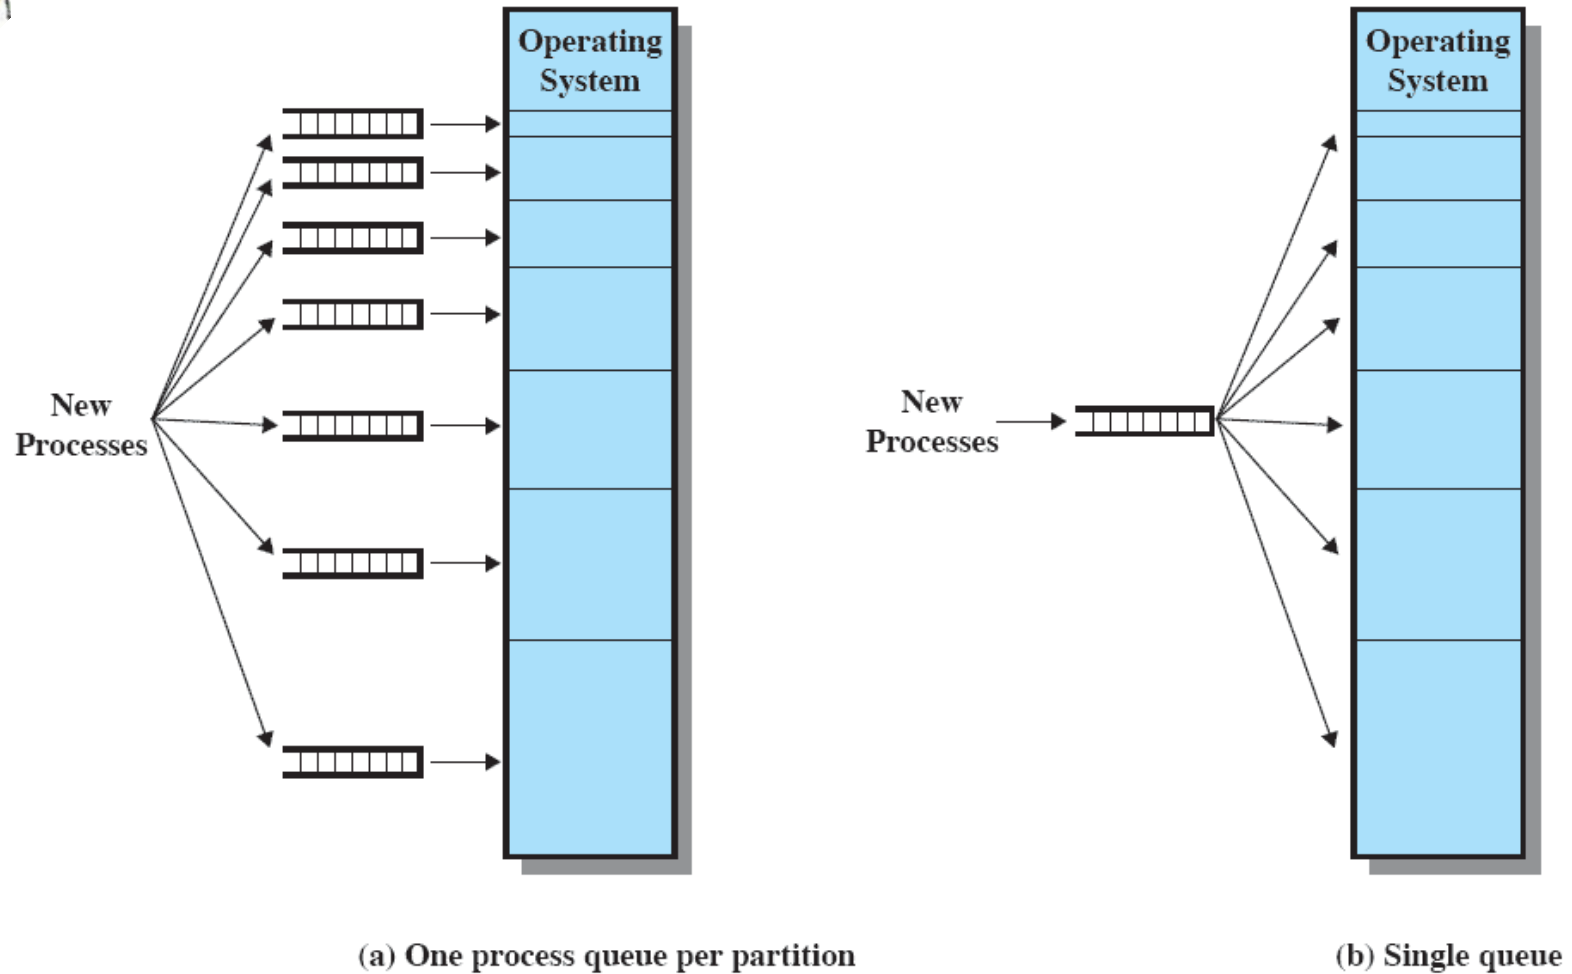





Figure 7.3 Memory Assignment for Fixed Partitioning



Remaining Problems with Fixed Partitions

- The number of active processes is limited by the system
 - I.E limited by the pre-determined number of partitions
- A large number of very small process will not use the space efficiently
 - In either fixed or variable length partition methods





Dynamic Partitioning


- Partitions are of variable length and number
- Process is allocated exactly as much memory as required




Dynamic Partitioning

Example

- ***External Fragmentation***
- Memory external to all processes is fragmented
- Can resolve using ***compaction***
 - OS moves processes so that they are contiguous
 - Time consuming and wastes CPU time




Refer to Figure 7.4



Dynamic Partitioning


- Operating system must decide which free block to allocate to a process
- Best-fit algorithm
 - Chooses the block that is closest in size to the request
 - Worst performer overall
 - Since smallest block is found for process, the smallest amount of fragmentation is left
 - Memory compaction must be done more often






Dynamic Partitioning


- First-fit algorithm
 - Scans memory from the beginning and chooses the first available block that is large enough
 - Fastest
 - May have many process loaded in the front end of memory that must be searched over when trying to find a free block





Dynamic Partitioning

- Next-fit
 - Scans memory from the location of the last placement
 - More often allocate a block of memory at the end of memory where the largest block is found
 - The largest block of memory is broken up into smaller blocks
 - Compaction is required to obtain a large block at the end of memory



Allocation







Figure 7.5 Example Memory Configuration before and after Allocation of 16-Mbyte Block



Buddy System

- Entire space available is treated as a single block of 2^U
- If a request of size s where $2^{U-1} < s \leq 2^U$
 - entire block is allocated
- Otherwise block is split into two equal buddies
 - Process continues until smallest block greater than or equal to s is generated





Example of Buddy System

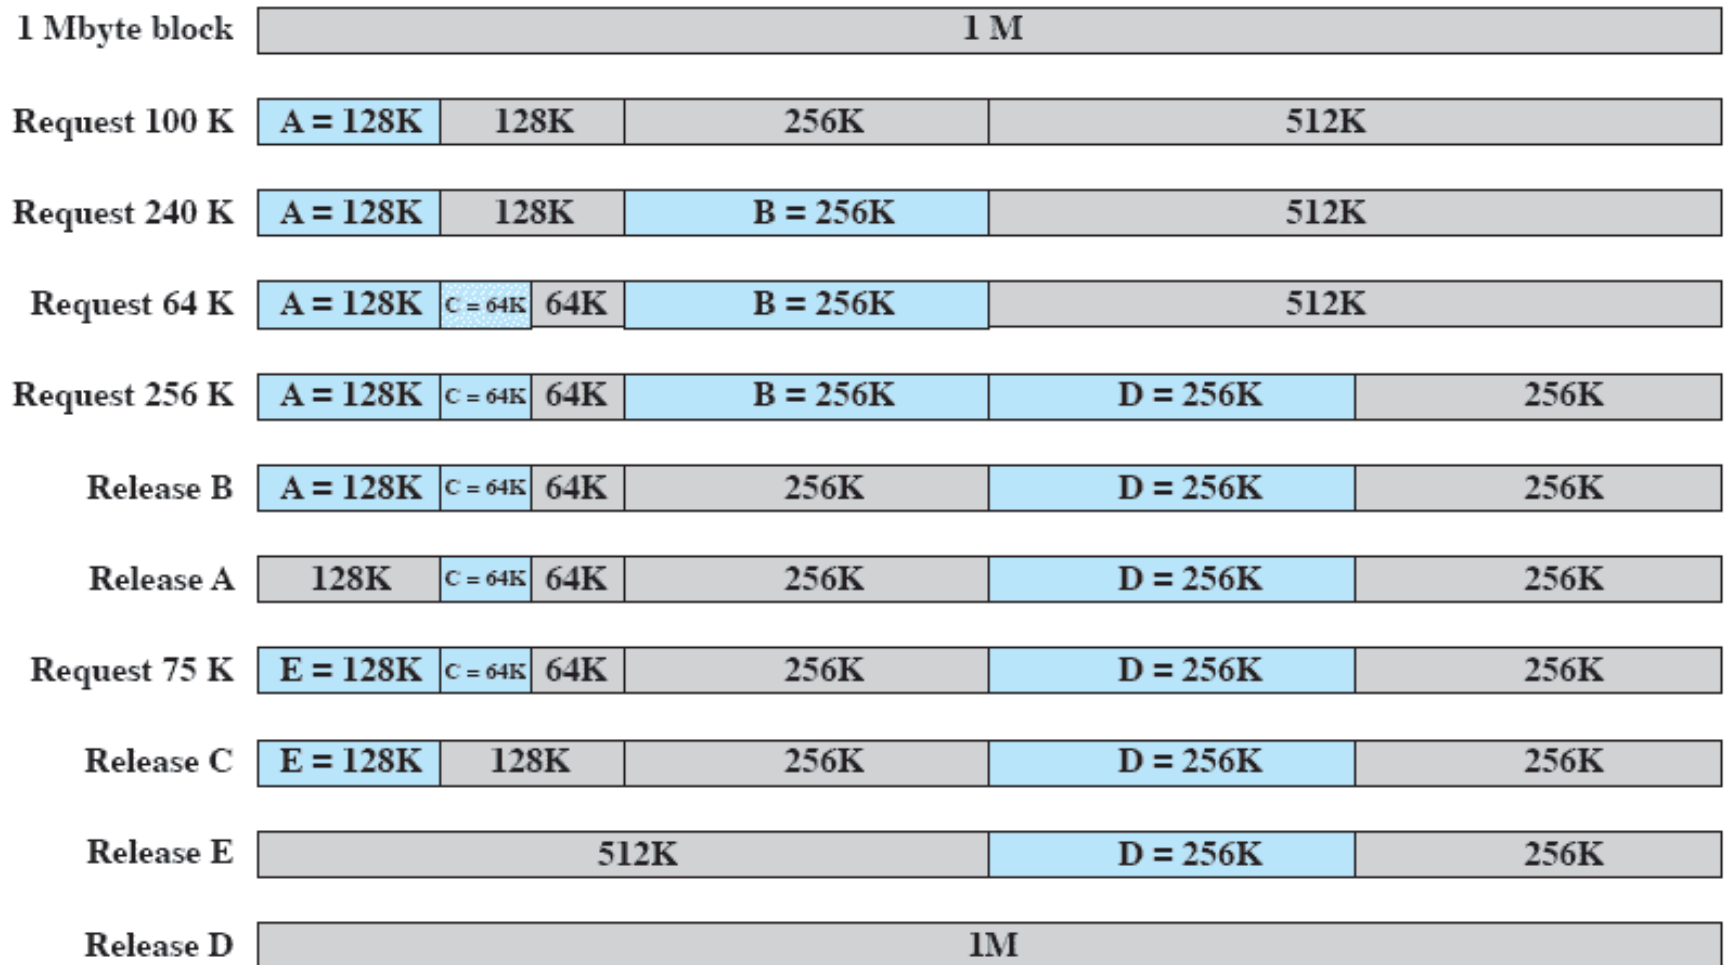


Figure 7.6 Example of Buddy System

Tree Representation of Buddy System

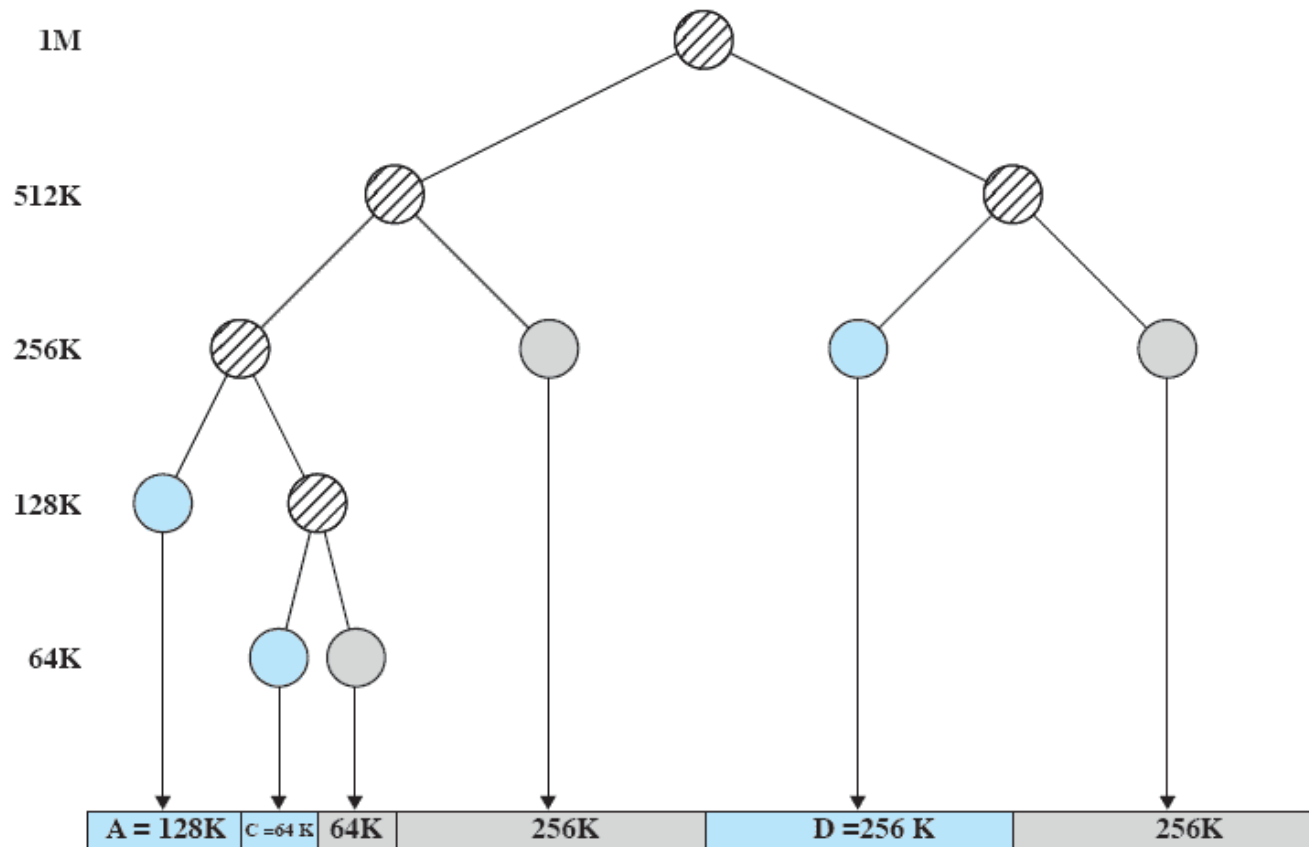




Figure 7.7 Tree Representation of Buddy System


Relocation

- When program loaded into memory the actual (absolute) memory locations are determined
- A process may occupy different partitions which means different absolute memory locations during execution
 - Swapping
 - Compaction



Addresses

- Logical
 - Reference to a memory location independent of the current assignment of data to memory.
- Relative
 - Address expressed as a location relative to some known point.
- Physical or Absolute
 - The absolute address or actual location in main memory.



Relocation

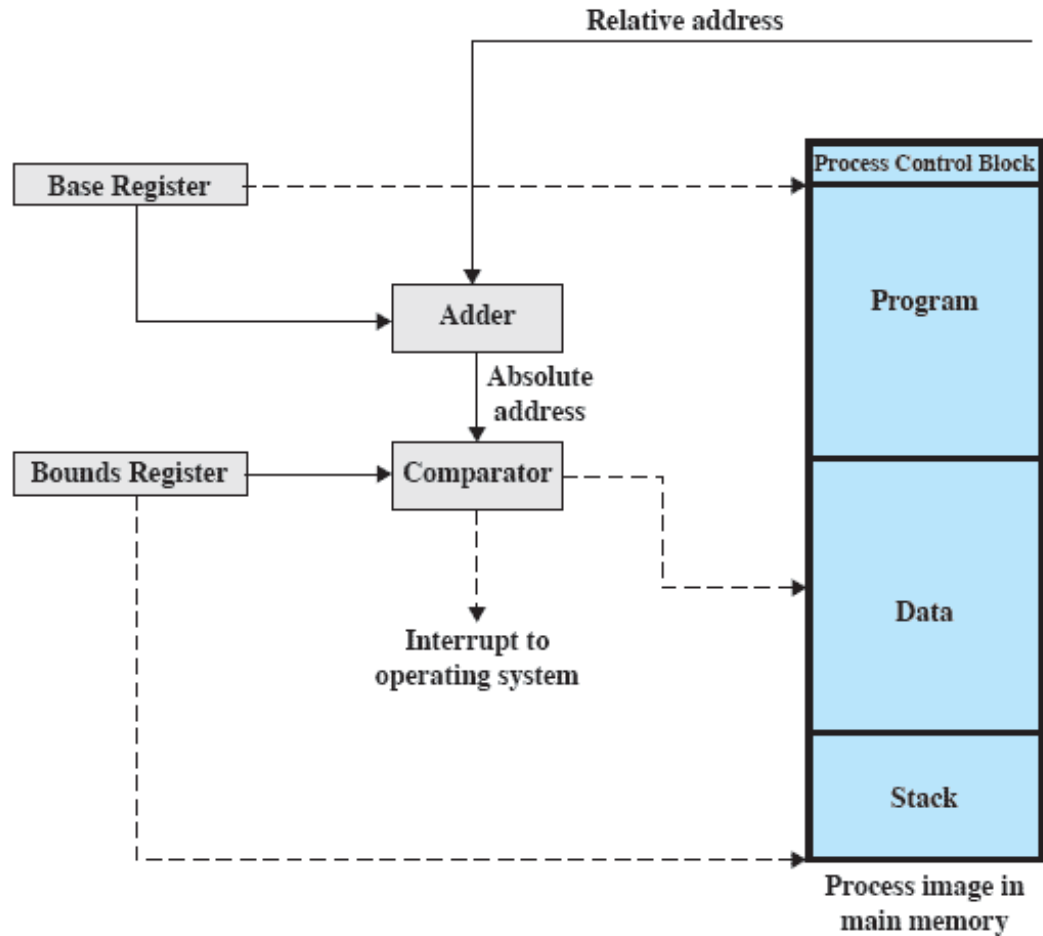





Figure 7.8 Hardware Support for Relocation



Registers Used during Execution

- Base register
 - Starting address for the process
- Bounds register
 - Ending location of the process
- These values are set when the process is loaded or when the process is swapped in






Registers Used during Execution


- The value of the base register is added to a relative address to produce an absolute address
- The resulting address is compared with the value in the bounds register
- If the address is not within bounds, an interrupt is generated to the operating system






Paging


- Partition memory into small equal fixed-size chunks and divide each process into the same size chunks
- The chunks of a process are called ***pages***
- The chunks of memory are called ***frames***






Paging

- Operating system maintains a page table for each process
 - Contains the frame location for each page in the process
 - Memory address consist of a page number and offset within the page





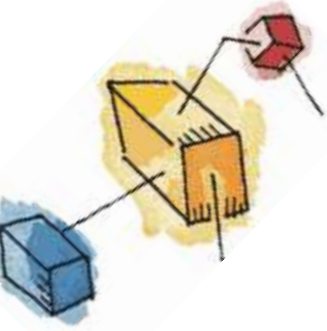
Processes and Frames



Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	



Page Table



0	0
1	1
2	2
3	3

Process A
page table

0	—
1	—
2	—

Process B
page table

0	7
1	8
2	9
3	10

Process C
page table

0	4
1	5
2	6
3	11
4	12

Process D
page table

13
14

Free frame
list

Figure 7.10 Data Structures for the Example of Figure 7.9 at Time Epoch (f)



Logical Addresses


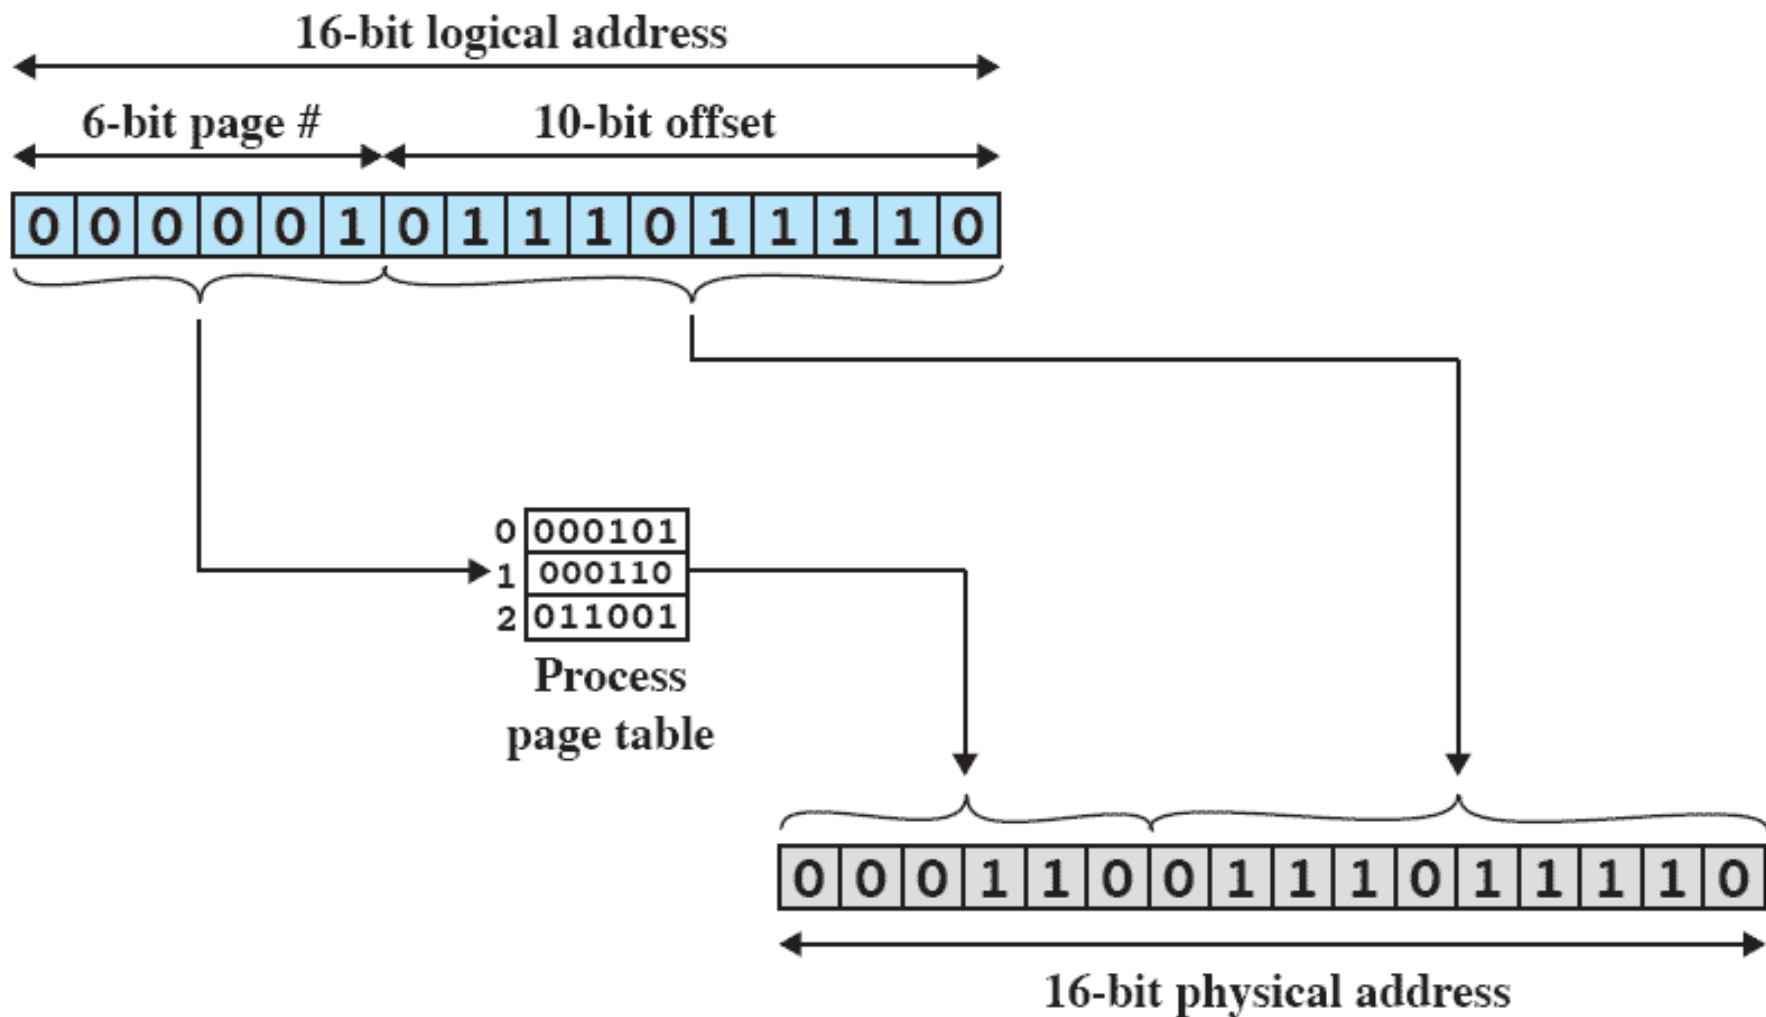



Figure 7.11 Logical Addresses

Paging




(a) Paging

Paging

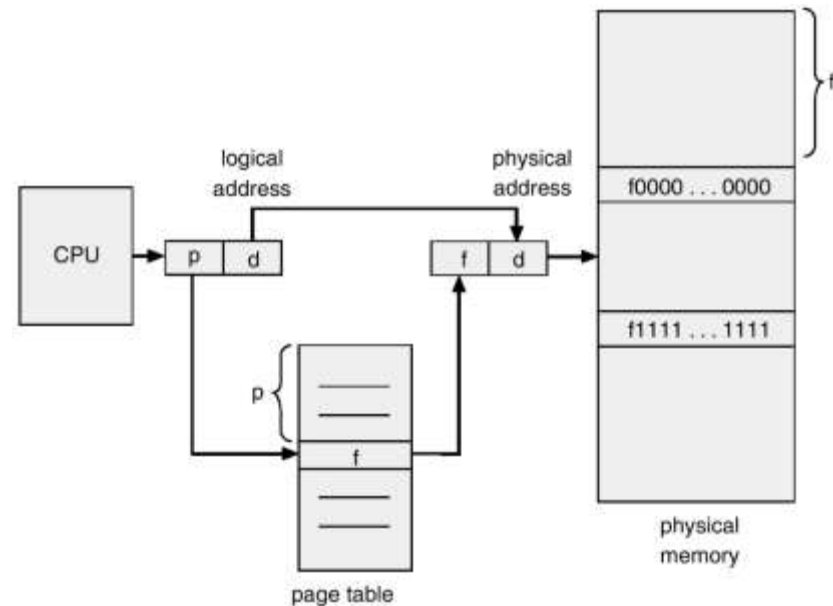


- Memory-management technique that permits the physical address space of a process to be non-contiguous.
- Each process is divided into a number of small, fixed-size partitions called pages or Divide logical memory into blocks of same size called **pages**.
- Physical memory is divided into a large number of small, fixed-size partitions called **frames**.
- When a process is to be executed, its pages are loaded into any available memory frames.
- ~~Some internal fragmentation , but no external fragmentation.~~




Address Translation Scheme

- Operating system maintains a page table for each process
 - Contains the frame number for each process page
 - To translate logical to physical addresses
- A logical address is divided into:
 - Page number (p) – used as an index into a page table which contains base address of each page in physical memory.
 - Page offset (d) – combined with base address to define the physical memory address that is sent to the memory.



- Paging Model of memory is shown:



Page 0
Page 1
Page 2
Page 3

Logical memory

0	1
1	4
2	3
3	7


Page table

↓ Frame number

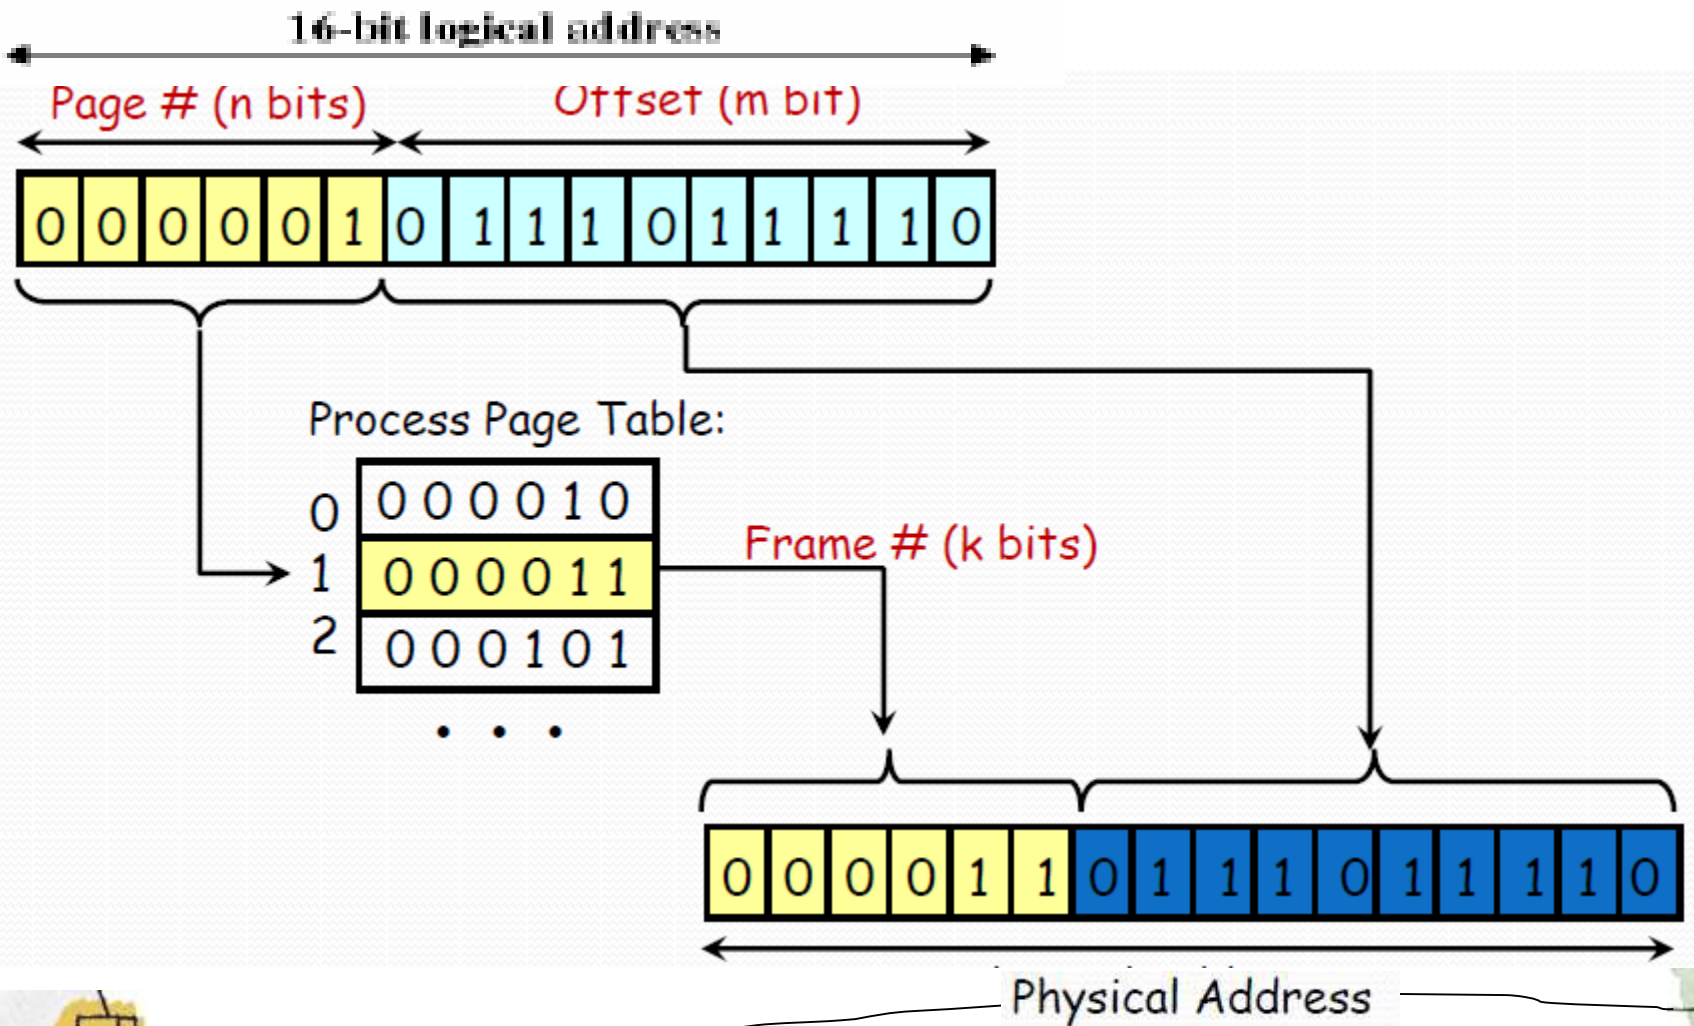
0	
1	Page 0
2	
3	Page 2
4	Page 1
5	
6	
7	Page 3

Physical memory

- To run a program of size n pages, need to find n free frames and load program.
- Set up a page table to translate logical to physical addresses




Address translation in paging






Segmentation

- A program can be subdivided into segments
 - Segments may vary in length
 - There is a maximum segment length
- Addressing consist of two parts
 - a segment number and
 - an offset
- Segmentation is similar to dynamic partitioning





Segmentation



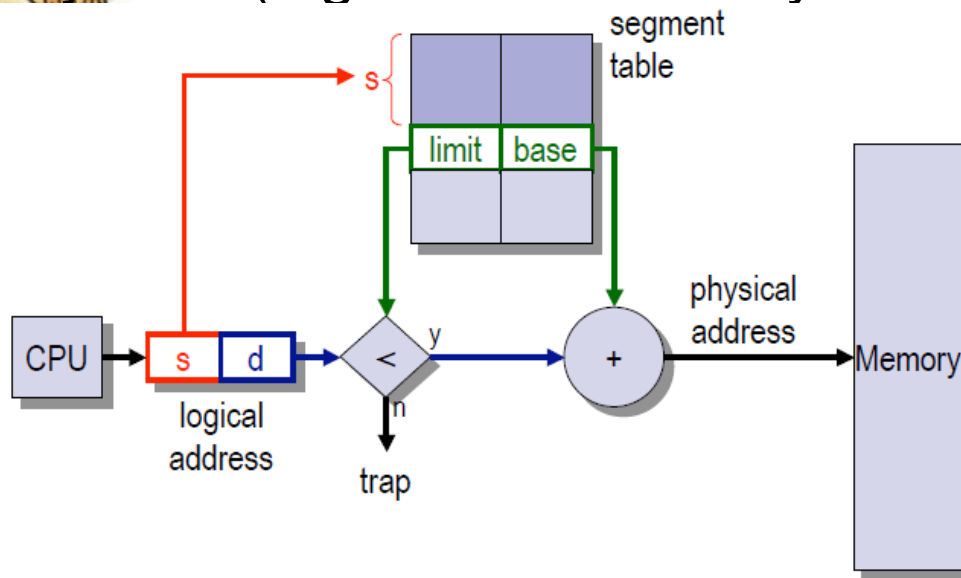
- A segment can be defined as a logical grouping of instructions, such as a subroutine, array, or a data area.
- A program is a collection of segments.
- Segmentation is a technique for managing these segments.
- Divide a process into unequal size blocks called segments.
- Each segment has a name and a length.
- Memory Management Scheme that supports user view of memory.



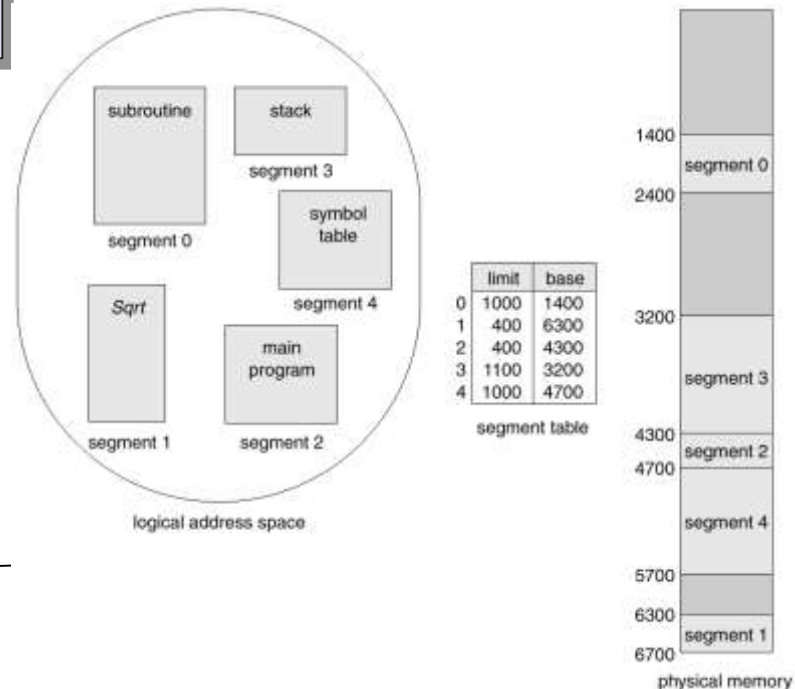
- 
- A segment is a logical unit such as
 - main program, procedure, function
 - local variables, global variables, common block
 - stack, symbol table, arrays
 - Protect each entity independently
 - Allow each segment to grow independently
 - Share each segment independently
-
- Each of these segments are of variable length and length will be defined in the program.



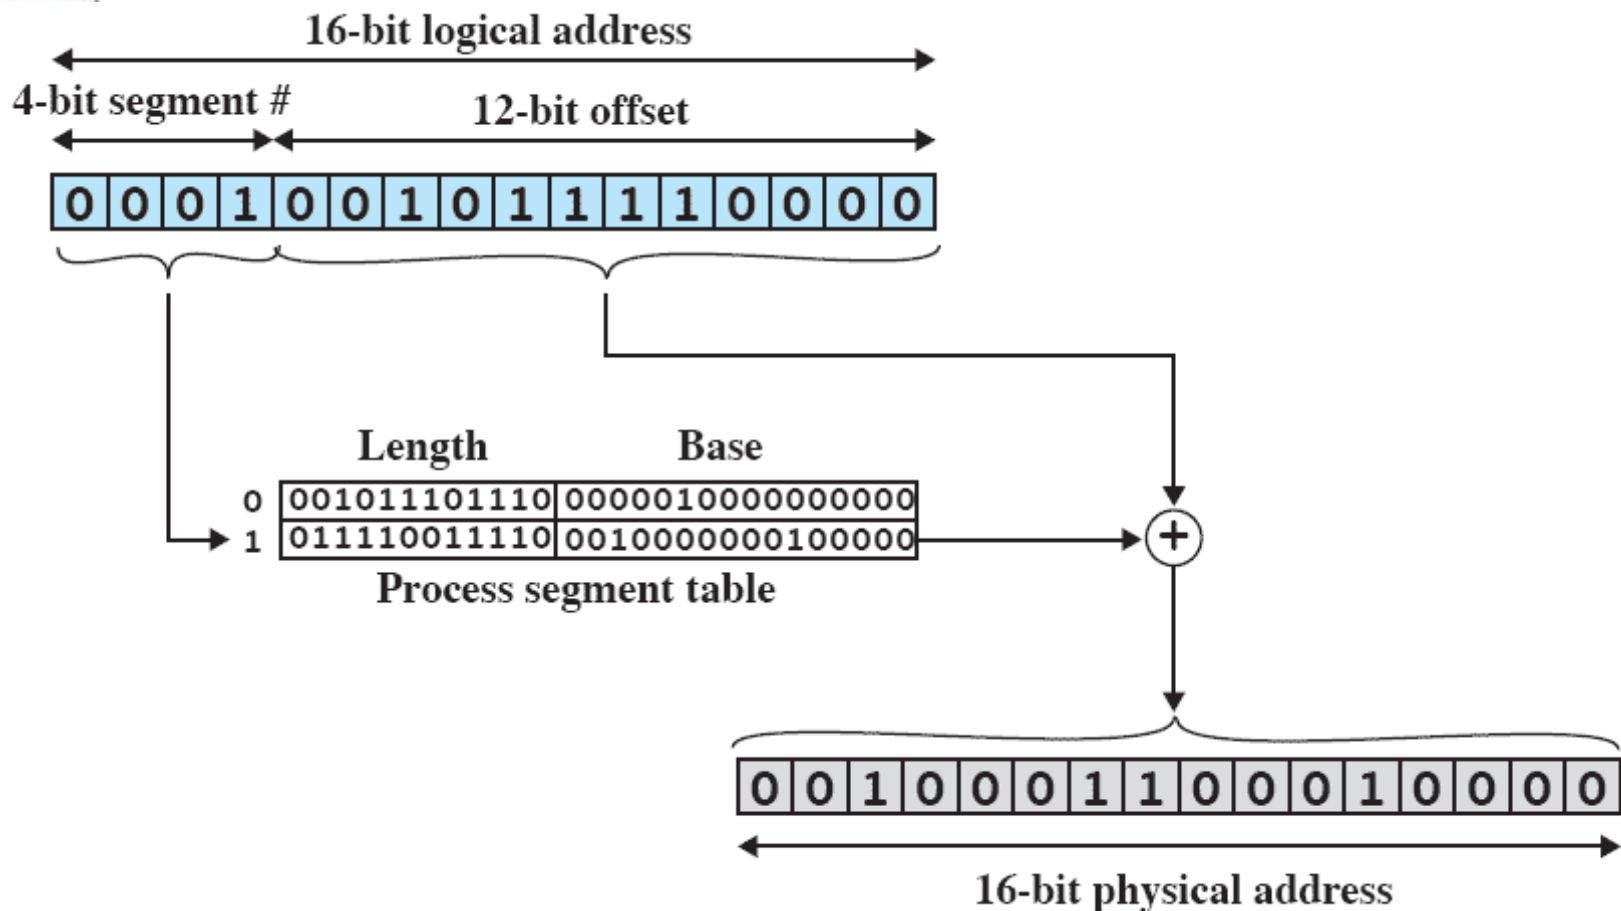
- If the offset is not within the limit, trap to the operating system.(logical address beyond end of segment).



- Segmentation Example:



Segmentation



(b) Segmentation

Figure 7.12 Examples of Logical-to-Physical Address Translation

Logical Addresses

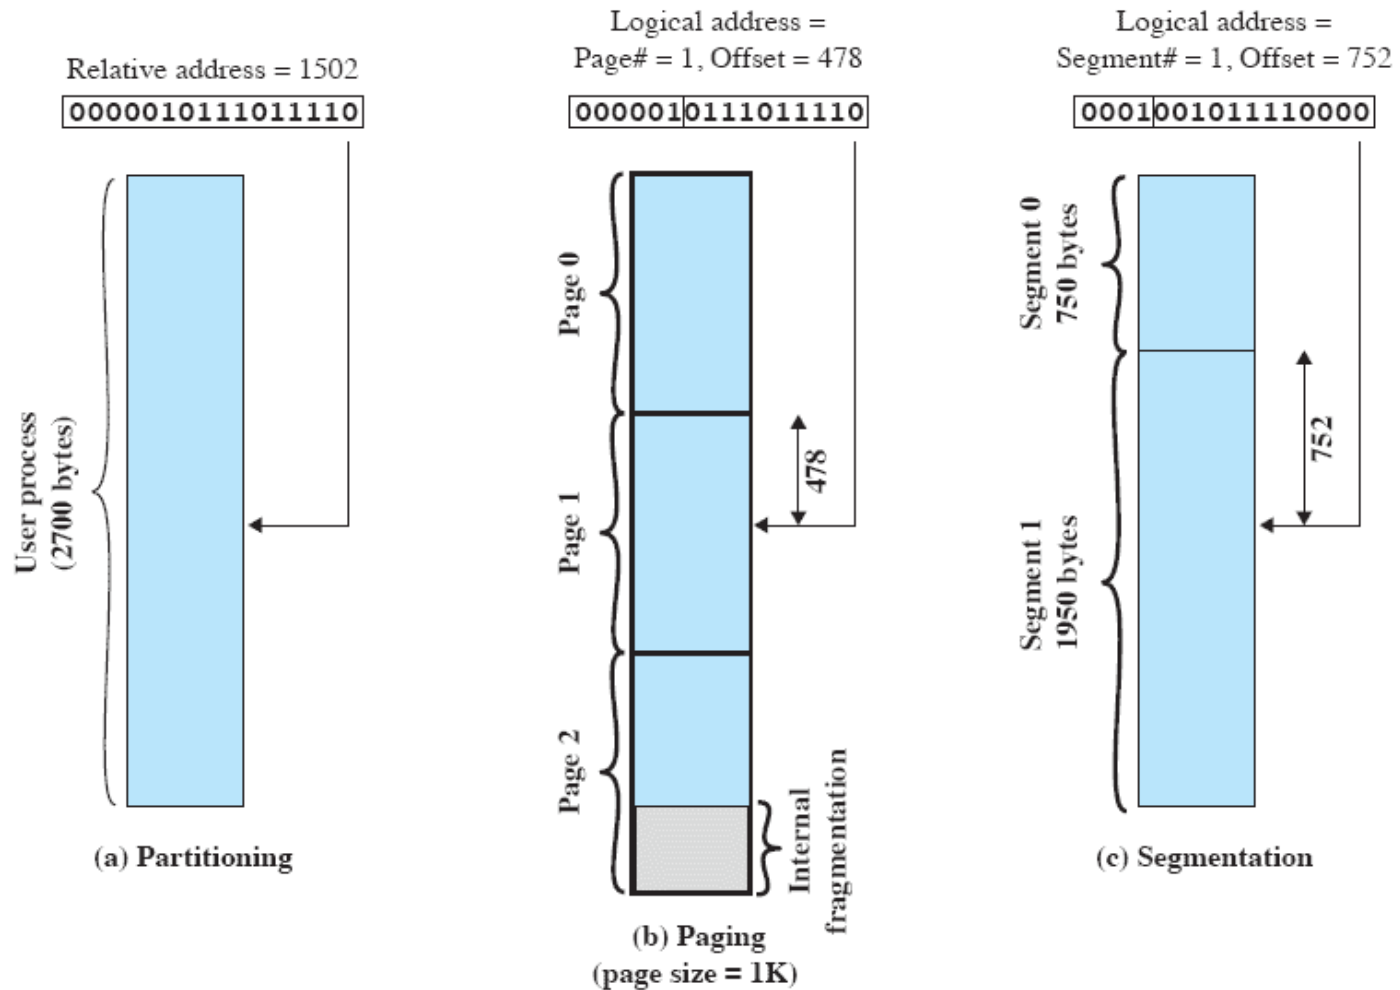


Figure 7.11 Logical Addresses



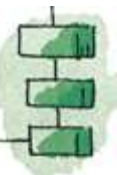
Advantages and Disadvantages

- Advantages:

- Each segment can be
 - located independently
 - separately protected
 - grow independently
- Segments can be shared between processes.
- Eliminate the fragmentation.

- Disadvantages:


- Allocation algorithms as for memory partitions
- External fragmentation, back to compaction problem.
 - Solution: combine segmentation and paging.






Segmentation vs. Paging

- Segment is good logical unit of information
 - sharing, protection
- Segmentation
 - Break process up into logical segments
- Segments are generally of different sizes
- Page is good physical unit of information
 - simple memory management
- Paging
 - Break process up into physical pages
 - Pages are the same size





Paging



The main memory partitioned into frames or blocks.

The logical address space divided into pages by compiler or MMU

Suffers from internal fragmentation or page break

Operating system maintains the page table.

Does not supports the user view of memory

Processor uses the page number and displacement to calculate absolute address.

Multilevel paging is possible

Segmentation

The main memory partitioned into segments.

The logical address space divided into segment specified by the programmer.


Suffers from external fragmentation.

Operating system maintains the segment table.

Supports the user view of memory

Processor uses the segment number and displacement to calculate absolute address.

Multilevel segmentation is possible. But no use.



Consideration

Paging

Segmentation

Need the programmer be aware that this technique is being used?	No	Yes
How many linear address spaces are there?	1	Many
Can the total address space exceed the size of physical memory?	Yes	Yes
Can procedures and data be distinguished and separately protected?	No	Yes
Can tables whose size fluctuates be accommodated easily?	No	Yes
Is sharing of procedures between users facilitated?	No	Yes
Why was this technique invented?	To get a large linear address space without having to buy more physical memory	To allow programs and data to be broken up into logically independent address spaces and to aid sharing and protection